

Name: _____

University of Illinois

Spring 2020

CS 446/ECE 449 Machine Learning
Homework 10: REINFORCE

Due on Tuesday May 5 2020, noon Central Time

1. [16 points] REINFORCE

We are given a utility $U(\theta) = \mathbb{E}_{p_\theta}[R(y)] = \sum_{y \in \mathcal{Y}} p_\theta(y) R(y)$ which is the expected value of the non-differentiable reward $R(y)$ defined over a discrete domain $y \in \mathcal{Y} = \{1, \dots, |\mathcal{Y}|\}$. Our goal is to learn the parameters θ of a probability distribution $p_\theta(y)$ so as to obtain a high utility (high expected reward), *i.e.*, we want to find $\theta^* = \arg \max_{\theta} U(\theta)$. To this end we define the probability distribution to read

$$p_\theta(y) = \frac{\exp F_\theta(y)}{\sum_{\hat{y} \in \mathcal{Y}} \exp F_\theta(\hat{y})}. \quad (1)$$

- (a) (3 points) If we are given an i.i.d. dataset $\mathcal{D} = \{(y)\}$ we can learn the parameters θ of a distribution via maximum likelihood, *i.e.*, by addressing

$$\max_{\theta} \sum_{y \in \mathcal{D}} \log p_\theta(y)$$

via gradient descent. State the cost function and its gradient when plugging the model specified in Eq. (1) into this program. When is this gradient zero?

Your answer:

- (b) (2 points) If we aren't given a dataset but if we are instead given a reward function $R(y)$ we search for the parameters θ by maximizing the utility $U(\theta)$, *i.e.*, the expected reward. Explain how we can approximate the utility by sampling from the probability distribution $p_\theta(y)$.

Your answer:

Name: _____

- (c) (3 points) Using general notation, what is the gradient of the utility $U(\theta)$ w.r.t. θ , *i.e.*, what is $\nabla_{\theta}U(\theta)$. How can we approximate this value by sampling from $p_{\theta}(y)$? Make sure that you stated the gradient in the form which ensures that computation via sampling from $p_{\theta}(y)$ is possible.

Your answer:

- (d) (5 points) Using the parametric probability distribution defined in Eq. (1), what is the approximated gradient of the utility? How is this gradient related to the result obtained in part (a)?

Your answer:

- (e) (3 points) In **A10_Reinforce.py** we compare the two forms of learning. Let the size of the domain $|\mathcal{Y}| = 6$, and let the groundtruth data distribution $p_{\text{GT}}(y) = 1/12$ for $y \in \{1, 6\}$, $p_{\text{GT}}(y) = 2/12$ for $y \in \{2, 5\}$, and $p_{\text{GT}}(y) = 3/12$ for $y \in \{3, 4\}$. The dataset \mathcal{D} contains $|\mathcal{D}| = 1000$ points sampled from this distribution. Further let $F_{\theta}(y) = [\theta]_y$, where $\theta \in \mathbb{R}^6$ and where $[a]_y$ returns the y -th entry of vector a . The reward function happens to equal the groundtruth distribution, *i.e.*, $R(y) = p_{\text{GT}}(y)$. What distribution p_{θ} is learned with the maximum likelihood approach? What distribution is learned with the REINFORCE approach? Explain why this is expected. Complete **A10_Reinforce.py** to answer these questions.

Your answer: